


A distributed solution of the distributed termination problem

Rana, S. P. "A distributed solution of the distributed termination problem." *Information Processing Letters* 17.1 (1983): 43-46.

Contents

- Introduction
- Rana's Algorithm
- Correctness
- Discussion



Introduction

Motivation

- **Distributed system**-> Lack of knowledge on the global state.
- A process has **no up-to-date knowledge** on the local states of other processes.
- Distributed **termination** and **deadlock** problems arise.
- **Termination**: several algorithms, but rely upon **designation of a process** to detect global termination.
- **Proposal**: **any** of the processes may initiate a wave to detect global termination.

Rana's algorithm

Motivation: an incorrect detection algorithm

Idea: detection & termination waves.

Basic messages are acknowledged.

A process is **quiet** if it is **passive** and all the **messages are acknowledged**.

1. Quiet process starts a wave tagged with its ID.
2. Only quiet processes take part of this wave.
3. If a wave completes, its initiator says “Termination”.

What is the problem?

Consider a process P that was not yet visited by a wave.

So P may take a quiet process Q active again by sending it a message!

Then P may become quiet and take part of the wave.

Eventually this wave may complete while q is active!

Solution: use a logical clock to provide each event with a time-stamp!

Assumptions (Rana's Algorithm)

Message-passing communication (control & basic), which can be time-stamped



Unique identification number 1,2,3...

Hamiltonian cycle 

Decentralized system



Symmetric processes

Synchronized processes (logical clock)



A process P_i can note down the clock when B_i is satisfied

Usual assumptions

General idea

A process is **passive** if B_i is satisfied, in other case it is **active**.


If an **active process** P_i satisfies B_i it sends a **time-stamped detection message + counter**.



If an **active process** receives a **detection message** it **purges it**.

If a **passive process** receives a **detection message**:

- If counter = n then **termination message**.
- If counter \neq n
 - If the message has a greater time-stamp then it **sends message** and **counter++**.
 - Otherwise it **purges it**.



The algorithm (CSP notation)

$P :: [P_1 \parallel \dots \parallel P_n]$

where, for each $1 \leq i \leq n$, $P_i :: * [S_i]$.

// Definimos el proceso concurrente
 P

$P_i :: B_i := \text{false};$

$* [S_i$

$\square B_i \rightarrow \text{BTIME}_i := \text{CLOCK-TIME}$

// $S_i = P_i$ waits for messages (if it receives a message then B_i becomes false)

BTIME_i = Time when P_i last satisfied B_i

CLOCK-TIME = current clock-time of process P_i


The algorithm (CSP notation)

```

Pi :: Bi := false;
• [ Si
  □ Bi → BTIMEi := CLOCK-TIME;
    TIME := BTIMEi;   COUNT := 1
    Pi+1 ! detection-message (TIME, COUNT)
  □ Pi-1 ? detection-message (TIME, COUNT) →
    [ ¬Bi → purge the message
    □ Bi →
      [ TIME < BTIMEi → purge the message
      □ TIME ≥ BTIMEi →
        COUNT := COUNT + 1;
        Pi+1 ! detection-message
          (TIME, COUNT)
      ]
    ]
]

```

If B_i becomes passive then
stamps and **sends detection**
message



P_i receives **detection message**:

If **active** it purges it.

If **passive** then:

If finished after TIME it purges message

If finished before TIME it joins
detection wave and **sends det.**
message


The algorithm (CSP notation)

```


Pi :: Bi := false;
* [ Si
  □ Bi → BTIMEi := CLOCK-TIME;
    TIME := BTIMEi;
    COUNT := 1;
    Pi+1 ! detection-message (TIME, COUNT)
  □ Pi-1 ? detection-message (TIME, COUNT) →
    [ COUNT = n →
      Pi+1 ! terminate-message;
      TERMINATE
    □ COUNT ≠ n →
      [ ¬Bi → purge the message
        □ Bi →
          [ TIME < BTIMEi → purge the message
            □ TIME ≥ BTIMEi →
              COUNT := COUNT + 1;
              Pi+1 ! detection-message
                (TIME, COUNT)
          ]
        ]
      ]
    ]
  Pi-1 ? termination-message →
    Pi+1 ! termination-message;
    TERMINATE
]

```

} If *COUNT*=*n*, P_i sends termination message and terminates



} If P_i receives termination message then sends message and terminates itself




Correctness

Correctness: assertion 1



If the global termination condition is satisfied, termination will be eventually detected



There is no possibility of detecting false termination

Discussion

Remarks

- Difficult to estimate # of messages passed in the **detection phase**.
- Exactly n messages are passed in the **termination phase**.
- **More than one** process can detect termination: no problem (other algorithms do have problems on this!)
- Other processes require to generate **spanning tree**.
- Limitation: synchronous communication (CSP), **messages on transit**. Approaches to deal with asynchronicity:
 - Modify global condition: $(B_i \text{ is true } \forall i) \text{ \& (no message on transit)}$
 - Modify local predicates B_i : quiet = passive & all messages are acknowledged.

References

- Fokkink, Wan. *Distributed Algorithms: An Intuitive Approach*. MIT Press, 2013.
- Lamport, Leslie. "Time, clocks, and the ordering of events in a distributed system." *Communications of the ACM* 21.7 (1978): 558-565.
- Hoare, Charles Antony Richard. *Communicating sequential processes*. Springer New York, 1978.
- Van Wezel, Michiel C., and Gerard Tel. "An assertional proof of Rana's algorithm." *Information processing letters* 49.5 (1994): 227-233.

Thanks