


Another advantage of **Free Choice**: Completely asynchronous agreement protocols

Ben-Or, Michael. "Another advantage of free choice (extended abstract): Completely asynchronous agreement protocols." *Proceedings of the second annual ACM symposium on Principles of distributed computing*. ACM, 1983.

Contents

- Introduction
- The consensus problem
- A consensus protocol
- Byzantine agreement
- Efficiency



Introduction

Context

- **Consensus:** termination, agreement and non-triviality
- **Fischer, Lynch and Paterson (FLP):** no consensus in asynchronous systems with a single undetectable fail.
- **Model extensions:** randomization, timing assumptions, failure detectors, strong primitives.
- This paper provides the **first randomized consensus protocol** (not very efficient).
- **Idea:** allow non-terminating executions but with probability 0.

The consensus problem

Assumptions and problem description

- Completely **asynchronous** strongly **connected** system
- Agreement on a $\{0,1\}$ value
- A process **decides** by setting a value (cannot be changed).
- **Operations:** send(Q,m), receive(P)
- Message **buffer M** to keep undelivered messages
- **Configuration:** N processes states + contents of M
- **Step:** a single process takes one configuration to another.

Assumptions and problem description

We need to handle probabilities of executions.

In each state of the system there may be a large number of operations of different processes that may occur next.

antes: the choice has been implicit in the choice of a single execution, now we want to consider ensembles of many executions, where the probability of individual executions is determined by the return values of the COIN-FLIP operation.

- **Coin flip operations:** returns “random” value.
- **Strong adversary-“scheduler”:**
 - Controls message system and which process makes a step.
 - Function from partial executions to operations: it **chooses which operation is next**.
 - Adversary observes the **entire history**.
- **t-correct schedule:** in an infinite run with at most t processes make a finite number of steps (**t-failures**) and **any message is delivered** if the process makes an infinite number of steps.

A consensus protocol

Overview

Idea: An infinite repetition of rounds to exchange info

- In round r , process P only handles messages with timestamp r
- Each round has two phases: **voting** and **ratification**
- In the **voting**, each process sends a value which is a function of the values collected in the previous round (the first value is the input value)
- In the **ratification**, each process sends a value which is a function of the collected values

Algorithm

Input: boolean value input

Output: boolean value stored in output

Data: boolean preference, integer round

begin

 preference \leftarrow input

 round \leftarrow 1

while *true* **do**

 send (1, round, preference) to all processes

 wait to receive $n - t$ (1, round, *) messages

if *received more than $n/2$ (1, round, v) messages* **then**

 send (2, round, v, *ratify*) to all processes

else

 send (2, round, ?) to all processes

end

 wait to receive $n - t$ (2, round, *) messages

if *received a (2, round, v, ratify) message* **then**

 preference \leftarrow v

if *received more than t (2, round, v, ratify) messages* **then**

 output \leftarrow v

end

else

 preference \leftarrow CoinFlip()

end

 round \leftarrow round + 1

end

Step 1

Step 2: voting

Step 3: ratification

Step 4

Algorithm 1: Ben-Or's consensus protocol. Adapted from [20].

Consensus is guaranteed

Theorem 1. Let $N > 2t$. For any t -correct schedule and any initial values of the processes, the above protocol guarantees, with probability 1, that:

- (i) all the processes will eventually decide on the same value v
- (ii) if all processes start with the value v , then within one round they will all decide v
- (iii) if for some round r , some process decides v in step 3, then all other processes will decide v within the next round

Remark: If $N \leq 2t$ then consensus is impossible, since the schedule can then simulate a network partition.

Agreement is guaranteed

- At most one value can receive a majority of votes in the first stage of a round, so for any two messages $(2;r;v;\text{ratify})$ and $(2;r;v';\text{ratify})$, $v=v'$
- If some process sees $t+1$ $(2;r;v;\text{ratify})$ messages, then every process sees at least one $(2;r;v;\text{ratify})$ message.
- If every process sees a $(2;r;v;\text{ratify})$ message, every process votes for v in the rst stage of round $r+1$ and every process that has not already decided decides v in round $r+1$

Validity is guaranteed

- Validity follows by a similar argument; if all processes vote for the their
- common input
- v
- in round 1, then all processes send (2
- $;r;v;$
- ratify
-) and decide
- in the second stage of round 1

Termination is guaranteed

- At most one value can receive a majority of votes in the first stage of a round, so for any two messages $(2;r;v;ratify)$ and $(2;r;v';ratify)$, $v=v'$
- If some process sees $t+1$ $(2;r;v;ratify)$ messages, then every process sees at least one $(2;r;v;ratify)$ message.
- If every process sees a $(2;r;v;ratify)$ message, every process votes for v in the r st stage of round $r+1$ and every process that has not already decided decides v in round $r+1$

Byzantine agreement

Description

- Arbitrary behaviour of faulty processes
- A process can determine the originator of a message he received.
- Scheduler:
 - When each process will make a step
 - What the faulty processes do.

Byzantine Protocol

Process P: Initial value x_p .

step 0: set $r = 1$.

step 1: Send the message $(1, r, x_p)$ to all the processes.

step 2: Wait till messages of type $(1, r, *)$ are received from $N-t$ processes. If more than $(N+t)/2$ messages have the same value v , then send the message $(2, r, v, D)$ to all processes. Else send the message $(2, r, ?, D)$ to all processes.

step 3: Wait till messages of type $(2, r, *)$ arrive from $N-t$ processes.

(a) If there are at least $t+1$ D-messages $(2, r, v, D)$, then set $x_p := v$.

(b) If there are more than $(N + t)/2$ D-messages then decide v .

(c) Else set $x_p = 1$ or 0 each with probability $1/2$

step 4: Set $r = r + 1$ and go to step 1.

Consensus is guaranteed

Theorem 2. Let $N > 5t$. For any t -correct schedule and any initial values of the processes, the above protocol guarantees, with probability 1, that:

- (i) all the processes will eventually decide on the same value v
- (ii) if all processes start with the value v , then within one round they will all decide v ; and
- (iii) if for some round r , some process decides v in step 3, then all other processes will decide v within the next round.

Remark: We do not know whether $N > 5t$ is the best possible bound to reach distributed Byzantine agreement.

Efficiency

But termination is slow

- The probability that the algorithm terminates in any given round may be exponentially small as a function of the number of processes, requiring exponentially many rounds.

Theorem 3. If $t = O(N^{1/2})$ then the expected number of rounds to reach agreement in protocols 1 and 2 is constant, (i.e. does not depend on N).

- For deterministic protocols it is known that Byzantine agreement is impossible in less than $t+1$ rounds of exchange of information.

References

- Ben-Or, Michael. "Another advantage of free choice (extended abstract): Completely asynchronous agreement protocols." *Proceedings of the second annual ACM symposium on Principles of distributed computing*. ACM, 1983.
- Aspnes, James. "Randomized protocols for asynchronous consensus." *Distributed Computing* 16.2-3 (2003): 165-175.

Thanks