

# Eventually Consistent: Not What You Were Expecting?


María Fernanda Mora Alba  
Luis Manuel Román García

**“In an ideal world there would be only one consistency model: when an update is made all observers would see that update”**

- Werner Vogels

# Contents

- Introduction
- Defining Eventual Consistency
- Relaxed Consistency Properties
- Prediction
- Empirical Measurement
- Comparison
- Future Work & Conclusion




# Introduction

# Introduction

Distributed computing is on the rise. The reasons of their importance are several:

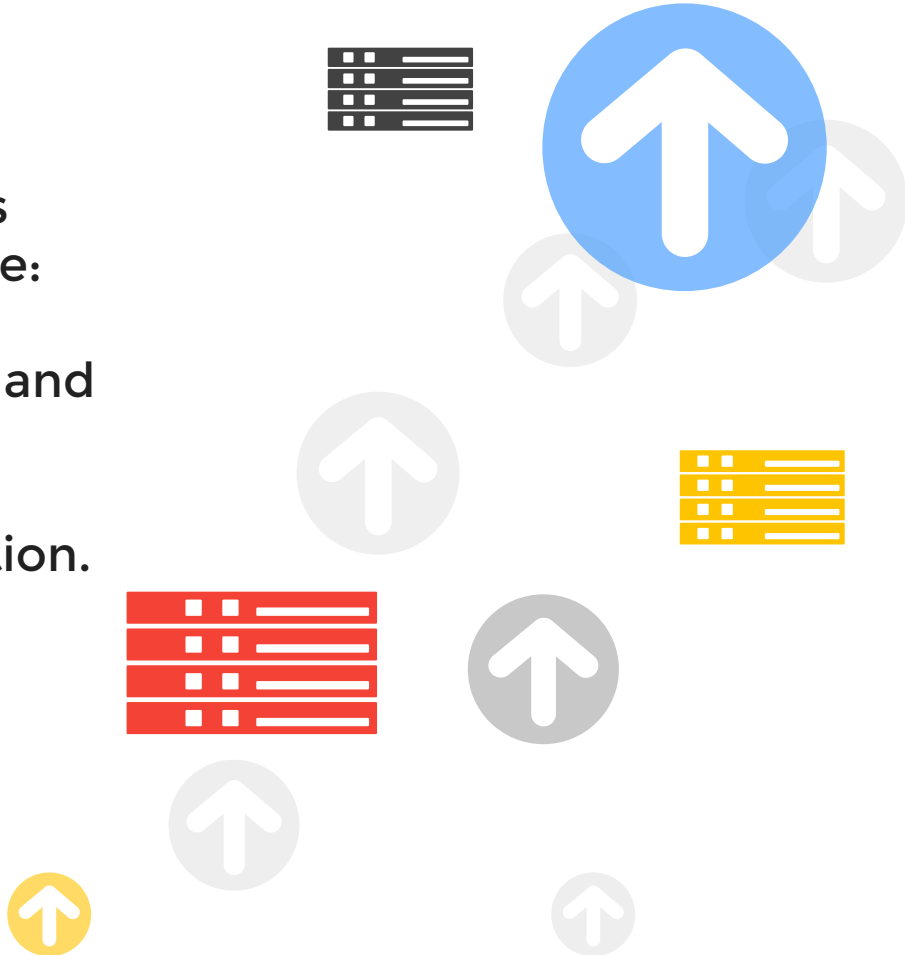
- Geographically distributed environment.
- Speed up.
- Resource sharing.
- Fault-tolerance.



# Introduction

Some of the main concerns regarding these systems are:


- Process independence and anonymity.
- Network topology.
- Degree of synchronization.
- Failures.



# Introduction

In most distributed database systems the concept of transaction is paramount. A **transaction** is a unit of consistent and reliable computation. Therefore, it must be:


- Atomic
- Consistent
- Isolated
- Durable



# Introduction

In the early days of distributed computing, **atomicity** was the main concern of the database designer.

This was mainly enforced by the principle posed by **Bruce Lindsay** in its **Notes on Distributed Databases** that required that the distribution of database systems should be **transparent**.






# Introduction

In the 1990's, with the grown of internet systems, this principles were revisited and **availability** became a major concern.

Then suddenly in a conference named **Principles of Distributed Computed** in 2000, **Eric Brewer** came out with a brilliant discovery...



# Introduction

## The curse of the CAP theorem


Consistency



Availability




Partition Resistance



# Defining Eventual Consistency

# Defining Eventual Consistency

It can be defined either as a property of the **underlying storage system** or as a behavior observed by a **client application**.




# Defining Eventual Consistency

It can be defined either as a property of the **underlying storage system** or as a behavior observed by a **client application**.

*“All replicas eventually receive all writes, and any two replicas that received the same set of writes have identical databases.”*

- Doug Terry

Storage system perspective




# Defining Eventual Consistency

It can be defined either as a property of the **underlying storage system** or as a behavior observed by a **client application**.

*“The storage system guarantees that if no new updates are made to the object, eventually all accesses will return the last updated value.”*

- Werner Vogels


Client application perspective



# Defining Eventual Consistency

Abstract definitions of eventual consistency leave open a number of questions regarding behaviour under **concurrent accesses** & **failure prone environments**.

It also leaves unresolved the issues regarding propagation speed and consistency among replicas.




# Defining Eventual Consistency

Due to the fact that eventually consistent systems are difficult to grasp, some other properties have been proposed, such as:

- Monotonic reads.
- Read my writes.
- Causal consistency.

In order to improve our understanding of such systems. In the following we are going to dig deeper in the several properties of an **eventually consistent systems**.






# Relaxed Consistency Properties

# Relaxed Consistency Properties

Comparing eventually consistent with fully sequential. **Linearizability**

How old is a value:

- Version-based staleness  
**k-atomicity**
- Time-based staleness  
 **$\Delta$ -atomicity**



# Prediction

# Prediction

**Probabilistically Bounded Staleness** (PBS) framework by Peter Bailis *et al.* to predict observed data staleness.

Estimates  $P(<k,t>\text{-staleness})$ .

When  $k=1$ , it estimates the probability of non-staleness.

Assumptions:

- Does not consider overlapping of reads and writes
- Does not model failures




# Empirical Measurement

# Empirical Measurement

**Difficult task:** concurrent operations  
difficult order identification, clients  
may always observe a different last  
updated value due to network partition

2 methodologies: **active measurement**  
and **passive analysis**.




# Active measurement

How eventual?: Measures **time** from the **write** to the **last read** that returns the **old value**.

Estimates the **convergence time** of the **replication protocol** (time needed to propagate new values to every replica).

**Challenges:** determine difference in time between writes and reads at different client nodes.

**YCSB++** (Yahoo! Cloud-serving benchmark) uses this approach.



# Active measurement

Can discover the range of update propagation times in an eventually consistent system: in **Amazon SimpleDB**, convergence of 90% of runs occurred in  $\leq 1$  sec and  $< 1\%$  of runs in  $> 4$  secs.

But doesn't indicate the % of the stale reads.

Frequent operations:  $\sim 100\%$  stale reads



Infrequent operations:  $\sim 0\%$  stale reads





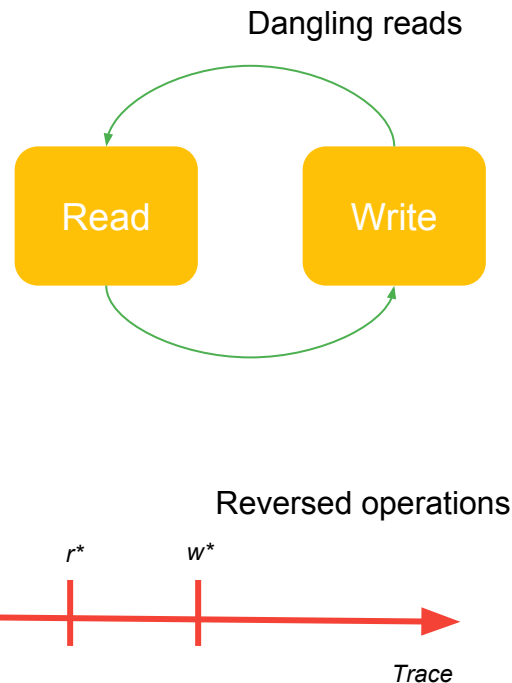
# Passive analysis

Prior consistency metrics:

1. Linearizability (binary property)
2. Delta-atomicity (bounds staleness of reads): **which delta** describes best the system? -> we **examine** operations in **trace**

**Challenges:** **collect the trace** i.e. for each operation  $O$  applied to  $X$  starting at  $t_1$  and finishing in  $t_2$  return the value of  $y$ .

**Problems:** **dangling reads** and **reversed operations**



# Comparison

# Comparison

## Address

Active measurement and PBS are **system-centric**: controlled workflow allows comparison between systems.

Passive analysis is **client-centric**: client interaction allows workload comparison.

## Model

Active measurement and PBS consider a **simplified model** (writes before reads).

Passive analysis considers a **general model** (allows overlapping).

## Cost vs taste

Passive analysis is **more expensive** but reflects better the **observed consistency** (stale reads).

# Future Work & Conclusion

# Future Work & Conclusion

## Future work

**Latency** impacts on user experience and revenue.

**Weak consistency** is used to **reduce latency**: how much is it reduced?

**Consistency & guarantees** awareness for differentiated pricing schemes: Amazon's DynamoDB, systems with user wishlist.

**Current challenges**: consistency verification (SLA, QoS): computation complexity of k-atomicity (only for  $k < 3$ ).

## Conclusion

**Eventual consistency** is not a binary property.

Characterizing, measuring and verifying it will **enrich services** and **improve user's experience**.

Thanks