



The byzantine generals problem

L. Lamport, R. Shostak, and M. Pease @ SRI International ACM Transactions on Programming Languages and Systems, July 1982, pp 382-401


Fernanda Mora
Luis Román

“A **distributed system** is one in which the **failure** of a computer you didn't even know **existed** can render your own computer **unusable**”

- Leslie Lamport

Contents


- Introduction
- Impossibility results
- A solution with oral messages
- A solution with signed messages
- Missing communication paths
- Reliable systems
- Conclusion & Discussion



Introduction

Consensus in synchronous faulty systems: BGP

- The problem of coping with **arbitrary, random failures** (byzantine) is the **Byzantine Generals Problem** (BGP)




We must
decide upon
a **common**
plan of
action

- Byzantine Army is divided in groups leadered by generals (nodes)
- Generals can communicate with each other using a messenger: **ATTACK or RETREAT**

Consensus in synchronous faulty systems: BGP

- The problem of coping with **arbitrary, random failures** (byzantine) is the **Byzantine Generals Problem** (BGP)



We must
decide upon
a **common**
plan of
action

Problem: some of the generals are
traitors (~**faulty nodes**)!

We don't know who the traitors are!

We want to guarantee:

1. All loyal generals decide upon the same plan of action.

They should use the same information $v(1), \dots, v(n)$

2. A small number of traitors cannot cause the loyal generals to adopt a bad plan.

We need a robust method: how does generals reach a decision?

We can have conditions on the *ith* general:

1. All loyal generals decide upon the same plan of action.

Any two loyal generals use the same value $v(i)$, for all i

2. A small number of traitors cannot cause the loyal generals to adopt a bad plan.

If the i th general is loyal, then the value he sends must be used by every loyal general as the value of $v(i)$

Byzantine Generals problem (BGP)

We can restrict on how a **single** general sends his value to others:

Formal BGP. A commanding general must send an order to his $n-1$ lieutenant generals such that:

IC1. **All loyal** lieutenants **obey** the **same** order.

IC2. If a **commander is loyal**, then **every loyal lieutenant obeys** the order he sends.

Byzantine Generals problem (BCP)

To **solve our original problem** (i.e. decide a plan), the *i*th general sends his value $v(i)$ by using a **solution to the BCP** to send the order “**use $v(i)$ as my value**”, with the other generals acting as the lieutenants.

Impossibility results

Impossibility of having $\frac{1}{3}$ or more traitors using oral messages

- 3 generals: 2 loyal, 1 traitor -> no solution!

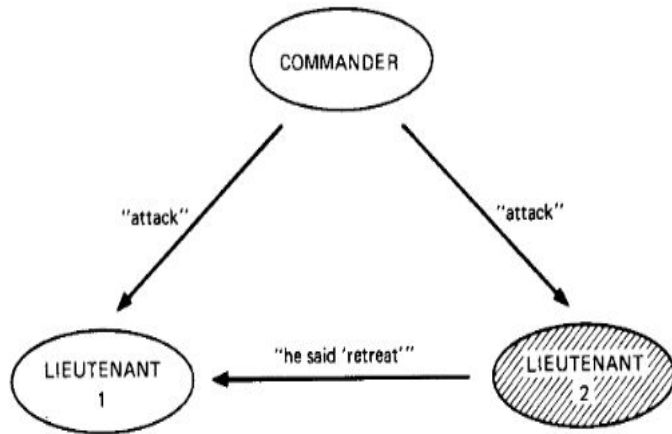


Fig. 1. Lieutenant 2 a traitor.

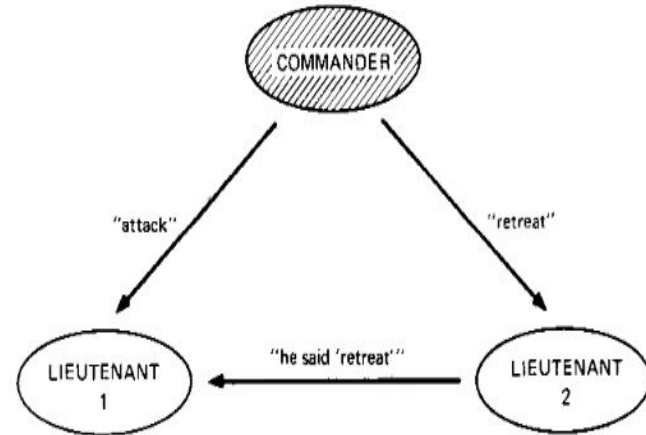


Fig. 2. The commander a traitor.

Impossibility of having $\frac{1}{3}$ or more traitors using oral messages

- **No solution** with fewer than $3m+1$ generals can cope with m traitors
- With **m traitors**, we need $n \geq 3m+1$ generals
- Reaching **approximate agreement** is as hard as exact agreement

A solution with oral
messages

Assumptions

1. Every **message** that is **sent** is **delivered** correctly.
2. The **receiver** of a message **knows who** sent it.
3. The **absence** of a message **can be detected**.

Prevent from
traitor interfering
communication

Prevents a
traitor's boycott

* m traitors and at least $3m+1$ generals

Oral message algorithm (recursive)

A commander sends an order to $n-1$ lieutenants a **majority function** such that

- $\text{majority}(v_1, \dots, v_{n-1}) = \text{mode}\{v_1, \dots, v_{n-1}\}$ or RETREAT if not order is received, or
- $\text{majority}(v_1, \dots, v_{n-1}) = \text{median}\{v_1, \dots, v_{n-1}\}$

Algorithm OM(0) (base case):

1. The commander sends its value to every $n-1$ lieutenants.
2. Each lieutenant uses the value he receives from the commander, or uses RETREAT if he receives no value.

Oral message algorithm (recursive)


Algorithm OM(m), $m > 0$ (recursive step):

1. The commander sends his value to every lieutenant.
2. For each i , let v_i the value lieutenant i receives from the commander, or else RETREAT. Lieutenant i acts as the commander in algorithm $OM(m-1)$ to send the value v_i to each of the $n-2$ other lieutenants.
3. $\forall i, j, i \neq j$, let v_i be the value lieutenant i received from lieutenant j in step 2 (using algorithm $OM(m-1)$), or else RETREAT. Lieutenant i uses the value $\text{majority}(v_1, v_2, \dots, v_n)$.


Oral message algorithm: remarks

- Lieutenants **recursively** forward orders to all the other lieutenants
- Algorithm $O(m-k)$ is called $(n-1)*\dots*(n-k)$ times to send a value prefixed with k lieutenants values
- Commander's order = majority (v_c, v_1, \dots, v_n)
- $v_i = \text{majority}(v_i, v_{i,2}, \dots, v_{i,n}), 1 \leq i \leq n$
- $v_{i,j} = \text{majority}(v_{i,j}, v_{i,j,3}, v_{i,j,4}, \dots)$
- **Unfolding:** $OM(m)$ invokes $n-1$ executions of $OM(m-1)$, which invokes $n-2$ executions of $OM(m-2)$, ...
- **Total number of messages:** $(n-1)*(n-2)*\dots*(n-m-1)$

Oral message algorithm: example



Algorithm OM(1): Lieutenant 3 a traitor



Algorithm OM(1): Commander a traitor

Two results

Lemma 1. For any m and k , Algorithm $OM(m)$ satisfies 1 if there are more than $2k+m$ generals and at most k traitors.

Theorem 1. For any m , algorithm $OM(m)$ satisfies conditions C1 and C2 if there are more than $3m$ generals, and at most m traitors.

A solution with signed
messages

More assumptions

1. A loyal general's signature cannot be forged.
2. Signatures can be authenticated.

Authentication requirements are compensated by a more resilient to faults algorithm

Algorithm SM(m) - overview

1. **Commander** sends a **signed** order to its lieutenants.
2. Each **lieutenant adds his signature** to that order and sends it to the others, who sign the order and send it to others, etc.
3. Each lieutenant **maintains a set of orders** he has received, i.e., the possible sets are:
 $\{\text{attack}\}, \{\text{retreat}\}, \{\text{attack}, \text{retreat}\}, \emptyset$
4. Lieutenant **takes action** according to the **value of the set**
 $\{\text{attack}, \text{wait}\}$ means the **commander is a traitor!**

Traitor

Remarks

- We need a function *choice* to choose an order from a set of orders V :
 - If $V=\{v\}$ then $\text{choice}(V)=v$
 - If $V=\{\emptyset\}$ then $\text{choice}(V)=\text{retreat}$
 - else $\text{choice}(V)=\text{median}(V)$ (for example)

Signed messages algorithm SM(m)

$V_i = \emptyset$, General 0 is the commander

1. Commander signs v and sends $v:0$ to all lieutenants.
2. For each lieutenant i :
 - a. If i receives $v:0$ and $V_i = \emptyset$
 - i. $V_i = \{v\}$
 - ii. sends $v:0:i$ to every other lieutenant.
 - b. If i receives $v:0:j_1 \dots j_k$ and $v \notin V_i$
 - i. Add v to V_i
 - ii. if $k < m$ sends $v:0:j_1 \dots j_k:i$ to all lieutenants $\notin \{j_1 \dots j_k\}$
3. When no more messages, i obeys order of **choice**(v_i)


Traitor

signature of m th
lieutenant is not
necessary

timeout
 $k \geq m$

Example: It's possible to have 3 generals and 1 traitor

Commander is traitor




Step 1.

Commander sends signed messages to L1 and L2

Example: It's possible to have 3 generals and 1 traitor

Commander is traitor



Step 2.

L1: $V_1 = \{\text{"attack"}\}$, sends "attack" : 0:1 to L2


L2: $V_2 = \{\text{"retreat"}\}$, sends "retreat" : 0:2 to L1

$V_1 = V_2 = \{\text{"attack"}, \text{"retreat"}\}$

L1, L2 know commander is traitor! (but no message is forged)

Example: It's possible to have 3 generals and 1 traitor

Commander is traitor




Step 3.

L1, L2 obey
choice ({“attack”, “retreat”})

IC1 and IC2 satisfied

Example: It's possible to have 3 generals and 1 traitor



Traitor

Lieutenant is traitor

Step 1.

Commander sends signed messages to L1 and L2

Step 2.

L2: $V_2 = \{\text{"attack"}\}$

L2 to L1 "retreat" : 0:2

L1: $V_1 = \{\text{"attack"}\}$

L1 receives "retreat" : 0:2

"L2 forged!"

L1 rejects the message

Step 3.

L1 "attacks"

IC2, IC1 satisfied

Some proofs

Theorem 2. For any m , algorithm $SM(m)$ solves the BGP if there are at most m traitors and $n \geq m+2$.

Sketch: Case 1: Commander is loyal: every loyal i receives v , and every other loyal receives v . Then $\forall_i v, IC1 \rightarrow IC2$.


Case 2: Commander is traitor: consider a loyal lieutenant j with signature list $|S|=m+1$. At least one i in S is loyal. The message accepted and sent by i must be accepted by every loyal lieutenant j , and viceversa. $V_i = V_j \rightarrow IC1$

- $SM(m)$ sends $(n-1)$ messages, each recipient $(n-2)$, ...
- $(n-1)(n-2) \cdots (n-m-1)$ messages to reach agreement
- Similar message complexity to $OM(m)$ but with better resilience.


Missing communication paths

Extensions

- **Network topology** can **restrict** communication.
- This makes Byzantine problem **more general**.
- For **simple graphs** we can extend OM(m), SM(m)
- OM(m,3m) solves BGP in a p-regular graph ($3m+1$ nodes at least) (is the same as OM(m))
- SM($n-2$) solves the GBP for n generals if loyal-graph is connected (can have missing links)



G is not 3-regular



G is 3-regular

Reliable *systems*

Ingredients

- **Majority voting** as a way to provide reliability
- What does we need?
 - **Input synchronization** of non-faulty processes (**IC1**)
 - If **input is non-faulty**, all non-faulty processes provide same output (**IC2**)
- **A1** – communication line vs node failure
 - No problem: OM(m) or SM(m) can deal with it
- **A2** – Fixed lines vs switching network
 - Not needed if we have A4
- **A3** – Timeouts
- **A4** – Cryptography

Conclusions & Discussion

Conclusions

- **Reliability** involves dealing with failure of components and is expensive: type of failure
- Byzantine failures produce arbitrary output making **agreement challenging**
- **BCP** used for input synchronization and handles m faults
- **Two solutions:** Oral and Signed Messages
- Expensive:

Time: message latencies and signatures

Messages: message paths $\leq m+1$ (optimal) , $O(n^{m+1})$ messages

Discussion

- Can we determine m ?
- How expensive is to implement a digital signature?
- How different is a “dumb” digital signature from an intelligent digital signature?
- If unfeasible to implement a digital signature, can we have $3m+1$ nodes?

Thanks